

Neural Network Training session

Afshin Sadeghi
Firas Kassawat

Based on The KG lab SS 2019



Table of Content

- PyTorch Installation
- PyTorch Tensors and Operations on Tensors
- Automatic Differentiation of Tensors
- Practice in class
- Implementing a linear regression using Pytorch
- Practice in class
- Exercise for home



PyTorch Installation

Operating System: Linux

Install miniconda from anaconda website:

<https://docs.conda.io/en/latest/miniconda.html>

Miniconda

| | Windows | Mac OS X | Linux |
|------------|------------------------|-------------------------|-------------------------|
| Python 3.7 | 64-bit (exe installer) | 64-bit (bash installer) | 64-bit (bash installer) |
| | 32-bit (exe installer) | 64-bit (.pkg installer) | 32-bit (bash installer) |
| Python 2.7 | 64-bit (exe installer) | 64-bit (bash installer) | 64-bit (bash installer) |
| | 32-bit (exe installer) | 64-bit (.pkg installer) | 32-bit (bash installer) |

Installation instructions



PyTorch Installation

Install pytorch using anaconda

| | | | | | | |
|-------------------|------------------------------------------------------|------------|------------|-------------------|------------|--------|
| PyTorch Build | Stable (1.0) | | | Preview (Nightly) | | |
| Your OS | Linux | | Mac | | Windows | |
| Package | Conda | | Pip | | LibTorch | Source |
| Language | Python 2.7 | Python 3.5 | Python 3.6 | | Python 3.7 | C++ |
| CUDA | 8.0 | | 9.0 | | 10.0 | None |
| Run this Command: | conda install pytorch-cpu torchvision-cpu -c pytorch | | | | | |



PyTorch Tensors and Operations on Tensors

A [torch.Tensor](#) is a multi-dimensional matrix containing elements of a single data type. In the tensor definition we can say it resides in CPU or GPU. It can be defined from a Python [list](#) or from a numpy array. (default is torch.FloatTensor or `float32`)

```
>>> torch.tensor([[1., -1.], [1., -1.]])
tensor([[ 1.0000, -1.0000],
        [ 1.0000, -1.0000]])

>>> torch.tensor(np.array([[1, 2, 3], [4, 5, 6]]))
tensor([[ 1,  2,  3],
        [ 4,  5,  6]])
```



PyTorch Tensors and Operations on Tensors

Defining the type of tensor:

```
>>> tensor = torch.ones((2,), dtype=torch.float64)
>>> tensor.new_full((3, 4), 3.141592)
tensor([[ 3.1416,  3.1416,  3.1416,  3.1416],
        [ 3.1416,  3.1416,  3.1416,  3.1416],
        [ 3.1416,  3.1416,  3.1416,  3.1416]], dtype=torch.float64)
```



PyTorch Tensors and Operations on Tensors

- Indexing and slicing notation:

```
>>> x = torch.tensor([[1, 2, 3], [4, 5, 6]])  
>>> print(x[1][2])  
tensor(6)
```

- For a tensor containing a single value: `item()`

```
>>> x = torch.tensor([[1]])  
>>> x  
tensor([[ 1]])  
>>> x.item()  
1
```



PyTorch Tensors and Operations on Tensors

- Creation can be by making a tensor of ones or zeroes:

```
>>> tensor = torch.tensor(), dtype=torch.float64)
>>> tensor.new_zeros((2, 3))
tensor([[ 0.,  0.,  0.],
        [ 0.,  0.,  0.]], dtype=torch.float64)
```

- Moving tensor to CPU or a core in GPU:

```
>>> cuda0 = torch.device('cuda:0')
>>> torch.ones([2, 4], dtype=torch.float64, device=cuda0)
tensor([[ 1.0000,  1.0000,  1.0000,  1.0000],
        [ 1.0000,  1.0000,  1.0000,  1.0000]], dtype=torch.float64, device='cuda:0')
```

`tensor.to('cpu')` (tensor here is the variable name)



PyTorch: Automatic Differentiation of Tensors

<https://pytorch.org/docs/stable/notes/autograd.html>

- By adding `.autograd` to tensor definition
- `Variable()`
- `nn.Parameter()`

$$\text{Tanh}(W_h * h^T + W_z * x^T)$$

$$\begin{cases} i2h = w_x \cdot x^T \\ h2h = w_h \cdot h^T \\ next_h = i2h + h2h \\ next_h = \tanh(next_h) \end{cases}$$

A graph is created on the fly

```
from torch.autograd import Variable
```

```
x = Variable(torch.randn(1, 10))  
prev_h = Variable(torch.randn(1, 20))  
W_h = Variable(torch.randn(20, 20))  
W_x = Variable(torch.randn(20, 10))
```





Practice in class

20 minutes:

Practice what you have learnt!

Use the Pytorch Tutorial website and do the followings:

Installation, tensors, operation on tensors, auto-diff, using cpu/gpu.

<https://bit.ly/2xABi3f>

<https://bit.ly/2UuEoRk>



PyTorch: Implementing a linear regression

torch.nn

`torch.nn.Module` : Base class for all neural network modules. For example:

`torch.nn.Conv2d` to define a convolutional layer

`torch.nn.Linear` to define a linear regression layer

`torch.nn.Embedding` to define lookup table that stores embeddings of a fixed dictionary and size.


Learning material online: <https://pytorch.org/docs/stable/nn.html?highlight=torch%20nn#module-torch.nn>

`torch.nn.Parameter` : A module that contains Tensor and is associated with a module class. When a Parameter is associated with a module as a model attribute, it is added to the parameter list automatically and can be accessed using the 'parameters' iterator of pytorch.



PyTorch: Implementing a linear regression

- **Creating and training linear regression model in PyTorch:**
-
- **Input: A set of N data samples in the form of (x,y)**
- - **x: input feature,**
 - **y: label/target.**
 -
- **Output: Trained linear model ($Wx + b$) to predict labels for unseen data**
- **Evaluation: Mean Square Error**



<https://hackernoon.com/linear-regression-in-x-minutes-using-pytorch-8eec49f6a0e2>



PyTorch: Defining a Network and learning the associated Classes

- **Creating Models in PyTorch:**
 - Create a Class
(LinearRegressionModel)
 - Declare your Forward Pass
 - Set the HyperParameters
(input_dim, output_dim)

Nn.linear (Ax+b)

```
class LinearRegressionModel(nn.Module):  
  
    def __init__(self, input_dim, output_dim):  
  
        super(LinearRegressionModel, self).__init__()  
        # Calling Super Class's constructor  
        self.linear = nn.Linear(input_dim, output_dim)  
        # nn.linear is defined in nn.Module  
  
    def forward(self, x):  
        # Here the forward pass is simply a linear function  
  
        out = self.linear(x)  
        return out
```

```
input_dim = 1  
output_dim = 1
```

<https://hackernoon.com/linear-regression-in-x-minutes-using-pytorch-8eec49f6a0e2>



PyTorch: Defining a Network and learning the associated Classes

Steps

1. Create instance of model
2. Select Loss Criterion
3. Choose Hyper Parameters

```
model = LinearRegressionModel(input_dim,output_dim)

criterion = nn.MSELoss()# Mean Squared Loss
l_rate = 0.01
optimiser = torch.optim.SGD(model.parameters(), lr = l_rate)
#Stochastic Gradient Descent

epochs = 2000
```

<https://hackernoon.com/linear-regression-in-x-minutes-using-pytorch-8eec49f6a0e2>



PyTorch: Defining a Network and learning the associated Classes

Training the Model in batch:

Zero_grad : clear grads

Forward pass (making prediction)

Criterion of optimizer(calculate error)

Back propagation(update weights)

Update step

```
for epoch in range(epochs):

    epoch +=1
    #increase the number of epochs by 1 every time

    inputs = Variable(torch.from_numpy(x_train))
    labels = Variable(torch.from_numpy(y_correct))

    #clear grads as discussed in prev post

    optimiser.zero_grad()

    #forward to get predicted values

    outputs = model.forward(inputs)
    loss = criterion(outputs, labels)
    loss.backward()# back props
    optimiser.step()# update the parameters
    print('epoch {}, loss {}'.format(epoch,loss.data[0]))
```



PyTorch: Defining a Network and learning the associated Classes

Finally, Print the Predicted Values

```
predicted
=model.forward(Variable(torch.from_numpy(x_train))).data.numpy()

plt.plot(x_train, y_correct, 'go', label = 'from data', alpha = .5)
plt.plot(x_train, predicted, label = 'prediction', alpha = 0.5)
plt.legend()
plt.show()
print(model.state_dict())
```

<https://hackernoon.com/linear-regression-in-x-minutes-using-pytorch-8eec49f6a0e2>



Practice in class

Students do LR in class

Questions?

Next Session?

LSTM, Relu, Concatenation and multiplication

Attention scores




PyTorch: Exercise for home 1

Create Neural Networks with (one, two and three) hidden layers and one output layer:

- Each node is a linear function($Wx+b$)
- `nn.Linear` can be a single node or a layer
- Hidden layer activation function
 - Tanh, sin, sigmoid, etc.
- Output node activation function
 - Linear function
- Dataset:
 - Experiment on 10 UCI regression dataset:
 - <https://archive.ics.uci.edu/ml/datasets.html>
- Loss function: MSE loss, Cross entropy loss
- Evaluation: Mean Square Error (Do experiments 20 times and report mean and variance of MSE)

Single Hidden Layer Neural Network with 4 Hidden Nodes






PyTorch: Exercise for home 2

Radial Basis Function Network with 4 Hidden Nodes

Create Radial Basis Function (RBF) Neural Network with one hidden layer and an output layer:

- Each node is a RBF node
- $g(x) = \exp(-b||x - m||)$
- Output node activation function
 - Linear function
- Dataset:
 - Experiment on 10 UCI regression dataset:
 - <https://archive.ics.uci.edu/ml/datasets.html>
- Loss function: MSE loss, Cross entropy loss
- Evaluation: Mean Square Error (Do experiments 20 times and report mean and variance of MSE)





References:

<https://docs.conda.io/en/latest/miniconda.html>

<https://pytorch.org/docs/stable/tensors.html#torch.Tensor>

<https://pytorch.org>

<https://hackernoon.com/linear-regression-in-x-minutes-using-pytorch-8eec49f6a0e2>

<https://blog.algorithmia.com/exploring-the-deep-learning-framework-pytorch/>